

Big Data Application Architecture - Final Project

ywang27@uchicago.edu

This document gives a brief description of my final project - A **Web Application for Querying and Submitting COVID-19 Data**. All the code for the project, including backend part and frontend part, are submitted along with this document. For simplicity, I select some part of code instead of all the lengthy code to show in this document.

Because I have very limited experience about node.js and web development before, I build my web page based on the flight-and-weather app. Although I cannot make many improvements to the code to make it rather different from the flight-and-weather app or more beautiful, I have tried my best in this part.

About

The Web Application for Querying and Submitting COVID-19 Data, as its name shown, is designed for express query of global COVID-19 data as well as submitting new covid-19 data occurring at the moment. All the data is collected from [COVID-19 Data Repository by the Center for Systems Science and Engineering \(CSSE\) at Johns Hopkins University](#). The application consists of two parts.

- Querying data. The application first accepts a country/region the user choosed (for example, United States). Then it shows the states/province of that country/region for further user choice (for example, Illinois). It also accepts a date value (for example, 2021-01-01). The application would show the accumulated data for this location from the beginning date of COVID-19 until the date user enters in, including
 - Confirmed cases
 - Deaths cases
 - Recovered cases
 - Active cases
 - Recovering rate (if applicable)
 - Death rate (if applicable)
- Submitting data. The application allows for submitting new COVID-19 data by providing a interface for user input. Similarly, a user input a new record by selecting country/region - state/province - date - case statistics. The added record would be stored and could be shown is it is queried.

Below are a couple of screeshots of the running application.

The page for querying the COVID-19 data:

[Click here to submit new covid-19 data](#)

Country/Region:

Please select a country/region

State/Province:

Please select country first

Date (YYYY-MM-DD):

2021-01-01

Submit

Select a country:

Da

Slovenia

Solomon Islands

Somalia

South Africa

South Sudan

Spain

Sri Lanka

Sudan

Summer Olympics 2020

Suriname

Sweden

Switzerland

Syria

Taiwan*

Tajikistan

Tanzania

Thailand

Timor-Leste

Togo

Tonga

Trinidad and Tobago

Tunisia

Turkey

US

Uganda

Ukraine

Select a state:

[Click here to submit new covid-19 data](#)

Country/Region:

US

State/Province:

✓ Please select country first

Alabama

Alaska

American Samoa

Arizona

Arkansas

California

Colorado

Connecticut

Delaware

Diamond Princess

District of Columbia

Florida

Georgia

Grand Princess

Guam

Hawaii

Idaho

Illinois

Indiana

Enter date "2021-01-01" and click submit:

[Click here to submit new covid-19 data](#)

Country/Region:

State/Province:

Date (YYYY-MM-DD):

Accumulated Covid-19 Data in US, Illinois until 2021-01-01

Confirmed	Active	Deaths	Recovered	Recovering Rate	Death Rate
963389	945411	17978	0	0%	1.87%

Click "Click here to submit new covid-19 data":

Submit Covid-19 Data

[Click here to query covid-19 data](#)

Country/Region:

State/Province:

Date:

Confirmed Deaths Recovered Active

Similarly, select the country and state, enter the required fields, and click submit:

Submit Covid-19 Data

Submit Covid-19 Data

[Click here to query covid-19 data](#)

Country/Region:

US

State/Province:

Illinois

Date:

2021-12-31

Confirmed Deaths Recovered Active

The return to the query page and query this new record:

[Click here to submit new covid-19 data](#)

Country/Region:

US

State/Province:

Illinois

Date (YYYY-MM-DD):

2021-12-31

Accumulated Covid-19 Data in US, Illinois until 2021-12-31

Confirmed	Active	Deaths	Recovered	Recovering Rate	Death Rate
10	8	1	1	10%	10%

which is the expected outcome.

Commands to Run

All the code are on the cluster `ec2-3-144-71-8.us-east-2.compute.amazonaws.com`, under the path `/home/hadoop/ywang27/final`. The files under `webapp` is also uploaded to `ec2-52-14-115-151.us-east-2.compute.amazonaws.com`, under the `/home/ywang27/final` where the app will be running on.

```
-final
  --BatchLayer
  --ServingLayer
  --SpeedLayer
    ---src
    ---target
  --webapp
    ---src
  --update_datalake_compute_batchview.sh
```

To run the application, first run the speed layer on `ec2-3-144-71-8.us-east-2.compute.amazonaws.com`

```
cd /home/hadoop/ywang27/final/SpeedLayer/target
spark-submit --master local[2] --driver-java-options "-
Dlog4j.configuration=file:///home/hadoop/ss.log4j.properties" --class StreamCovid uber-
SpeedLayer-1.0-SNAPSHOT.jar b-2.mpcs53014-kafka.198nfg.c7.kafka.us-east-
2.amazonaws.com:9092,b-1.mpcs53014-kafka.198nfg.c7.kafka.us-east-2.amazonaws.com:9092
```

The run the node.js on `ec2-52-14-115-151.us-east-2.compute.amazonaws.com` (here I use port 3027)

```
cd /home/ywang27/final/src
node app.js 3027 172.31.39.49 8070 b-1.mpcs53014-kafka.198nfg.c7.kafka.us-east-
2.amazonaws.com:9092,b-2.mpcs53014-kafka.198nfg.c7.kafka.us-east-2.amazonaws.com:9092
```

Go to `http://ec2-52-14-115-151.us-east-2.compute.amazonaws.com:3027/covid-show.html` and the application should be runing.

Where to find ...

- Tables in Hive: all the tables involved in my project are `ywang27_covid_all`, `ywang27_covid`, `ywang27_country_date`, `ywang27_country_date_for_hbase`
- Tables in Hbase: the Hbase table for my project is `ywang27_country_state_date`

Implementation

Datalake - HDFS

First download/update the original dataset from GitHub.

```
git pull
```

The original dataset is a bunch of csv files, eaching containing the data for one day. To create the datalake, I first concat all the cvs files as one file.

```
cat *.csv > all.csv
```

Then load it into **HDFS**

```
hdfs dfs -put all.csv /ywang27
```

Batch Layer - Hive

In **Hive**, create a Hive table `ywang27_covid_all` to store all the data in `all.csv`

```
CREATE EXTERNAL TABLE IF NOT EXISTS `ywang27_covid_all`(  
    # fields  
)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' NULL DEFINED AS ''  
STORED AS TEXTFILE;  
  
LOAD DATA INPATH '/ywang27/all.csv' OVERWRITE INTO TABLE ywang27_covid_all;
```

Filter out the record containing invalid date and save the results to a new table `ywang27_covid`

```
INSERT overwrite TABLE ywang27_covid  
SELECT * FROM ywang27_covid_all n  
WHERE n.last_update IS NOT NULL;
```

Create a new table `ywang27_country_date` to save batch view, which contains processed `province_state` and `update_date` fields

```
SELECT DISTINCT  
    if(province_state is not NULL, province_state, country_region) as province_state,  
    to_date(last_update) as update_date,  
    # ...  
FROM ywang27_covid
```


All the code for serving layer are in `final/BatchLayer`

Serving Layer - HBase

In **Hbase**, create a table `ywang27_country_state_date` with column family `data`

```
create 'ywang27_country_state_date', 'data'
```

In Hive, create a table `ywang27_country_date_for_hbase`, connect it with `ywang27_country_state_date`, and insert overwrite it by selecting from `ywang27_country_date`

```
CREATE EXTERNAL TABLE IF NOT EXISTS ywang27_country_date_for_hbase (  
    # fields  
)  
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
WITH SERDEPROPERTIES ('hbase.columns.mapping' =  
    ':key,data:confirmed,data:deaths,data:recovered,data:active')  
TBLPROPERTIES ('hbase.table.name' = 'ywang27_country_state_date');  
  
INSERT OVERWRITE TABLE ywang27_country_date_for_hbase  
SELECT CONCAT(country_region, province_state, update_date),  
    # other fields  
FROM ywang27_country_date;
```

All the code for serving layer are in `final/ServingLayer`.

The procedure for updating the datalake and batch views could be done by simply running `final/update_datalake_compute_batchview.sh`.

Speed Layer - Spark, Scala, Kafka

First create a topic `ywang27_covid` in **Kafka** for this project

```
kafka-topics.sh --create --zookeeper z-3.mpcs53014-kafka.198nfg.c7.kafka.us-east-  
2.amazonaws.com:2181,z-2.mpcs53014-kafka.198nfg.c7.kafka.us-east-2.amazonaws.com:2181,z-  
1.mpcs53014-kafka.198nfg.c7.kafka.us-east-2.amazonaws.com:2181 --replication-factor 1 --  
partitions 1 --topic ywang27_covid
```

Then create a **Scala** class `KafkaCovidRecord` for messages in Kafka in `KafkaCovidRecord.scala`

```
case class KafkaCovidRecord(  
  country: String,  
  state: String,  
  date: String,  
  confirmed: String,  
  deaths: String,  
  recovered: String,  
  active: String  
)
```

Then create the object `StreamCovid` in `StreamCovid.scala`.

In this object, create the hbase connection and get connected with `ywang27_country_state_year`

```
val hbaseConnection = ConnectionFactory.createConnection(hbaseConf)  
val covid_data = hbaseConnection.getTable(TableNames.valueOf("ywang27_country_state_date"))
```

Declare the **SparkStreaming** context with interval 5 seconds

```
val ssc = new StreamingContext(sparkConf, Seconds(5))
```

Set Kafka topic to `ywang27_covid`

```
val topicsSet = Set("ywang27_covid")
```

After we get the serialized records in kafka, read its value using the `KafkaCovidRecord`

```
val kfrs = serializedRecords.map(rec => mapper.readValue(rec, classOf[KafkaCovidRecord]))
```

Finally add the value to Hbase via `org.apache.hadoop.hbase.client.Put`

```

val processed_covid = kfrs.map(addToHbase)

def addToHbase(kfr: KafkaCovidRecord): String = {
  val rowkey = kfr.country + kfr.state + kfr.date
  val put = new Put(Bytes.toBytes(rowkey))
  put.addColumn(Bytes.toBytes("data"), Bytes.toBytes("confirmed"),
Bytes.toBytes(kfr.confirmed))
  put.addColumn(Bytes.toBytes("data"), Bytes.toBytes("active"),
Bytes.toBytes(kfr.active))
  put.addColumn(Bytes.toBytes("data"), Bytes.toBytes("deaths"),
Bytes.toBytes(kfr.deaths))
  put.addColumn(Bytes.toBytes("data"), Bytes.toBytes("recovered"),
Bytes.toBytes(kfr.recovered))
  covid_data.put(put)
  return "Updated speed layer"
}

```

Web - Node.js

The frontend part is quite similiar to flight-and-weather app.

In `app.js`, set the kafka topic to send to as `ywang27_covid`

```

kafkaProducer.send([ { topic: 'ywang27_covid', messages: JSON.stringify(report)}],
  function (err, data) {
    console.log(report);
    res.redirect('covid-submit.html');
  });

```

In order to implement that the second selection option lists (state list) changes as the first selection option list (country list), in function `setstate()` get the id of the second select tag, concatenate the all states/province according to the country value, and assign the concatenated string to the second select tag - `innerHTML` property. The function is called when the first selection (country) has been made.

```

function setstate(){
  let obj = document.getElementById("country");
  let country = obj.options[obj.selectedIndex].value;
  let states = country_state[country];
  let numofstates = states.length;
  let output = '<option value="">Please select country first</option>'

  for (var i = 0; i < numofstates; i++){
    output += '<option value="${states[i]}">${states[i]}</option>'
  }

  let $id_state = document.getElementById("state")

```

```
    id_state = document.getElementById( state );  
    $id_state.innerHTML = output  
}
```